

```
In [10]: 1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4 from torch.utils.data import Dataset, DataLoader
5 from collections import Counter
6 import torch.nn.functional as F
7 from sklearn.model_selection import train_test_split
8 import jieba
9 import numpy as np
10 import re
```

```
In [11]: 1 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
2 print("Using device:", device)
```

Using device: cuda

```
In [33]: 1 stopwords_file = "./cn_stopwords.txt"
2 with open(stopwords_file, "r", encoding="utf-8") as file:
3     stop_words_list = file.readlines()
4 stop_words_list = [line.strip() for line in stop_words_list]
```

```
In [19]: 1 with open("./骆驼祥子.txt", "r", encoding="utf-8") as file:
2     text = file.read()
3
4 text = re.sub(r'[\w\s]', '', text)
5 text = re.sub(r'\s+', ' ', text)
6 text = text.replace(" ", "")
7
8 raw_text_0 = list(jieba.cut(text, cut_all=False))
9 raw_text = [word for word in raw_text_0 if word not in stop_words_list]
10
11 vocab = set(raw_text)
```

```
In [20]: 1 CONTEXT_SIZE = 2 # 2 words to the left, 2 to the right
2 EMBEDDING_DIM = 50
3
4 # By deriving a set from `raw_text`, we deduplicate the array
5 vocab_size = len(vocab)
6
7 word_to_ix = {word: i for i, word in enumerate(vocab)}
8 data = []
9 for i in range(CONTEXT_SIZE, len(raw_text) - CONTEXT_SIZE):
10     context = (
11         [raw_text[i - j - 1] for j in range(CONTEXT_SIZE)]
12         + [raw_text[i + j + 1] for j in range(CONTEXT_SIZE)]
13     )
14     target = raw_text[i]
15     data.append((context, target))
16 print(data[:5])
```

```
[(['祥子', '骆驼', '祥子', '骆驼'], '介绍'), (['介绍', '祥子', '骆驼', '骆驼'],
'祥子'), (['祥子', '介绍', '骆驼', '外号'], '骆驼'), (['骆驼', '祥子', '外号',
'先'], '骆驼'), (['骆驼', '骆驼', '先', '说祥子'], '外号')]
```

```
In [21]: 1 class CBOW(nn.Module):
2
3     def __init__(self, vocab_size, embedding_dim, context_size):
4         super(CBOW, self).__init__()
5         self.embeddings = nn.Embedding(vocab_size, embedding_dim).to(device)
6         self.linear1 = nn.Linear(embedding_dim, 128).to(device)
7         self.linear2 = nn.Linear(128, vocab_size).to(device)
8
9     def forward(self, inputs):
10         embeds = torch.mean(self.embeddings(inputs), dim=0).view((1, -1))
11         out = F.relu(self.linear1(embeds))
12         out = self.linear2(out)
13         log_probs = F.log_softmax(out, dim=1)
14         return log_probs
15
16 losses = []
17 loss_function = nn.NLLLoss()
18 model = CBOW(len(vocab), EMBEDDING_DIM, CONTEXT_SIZE).to(device)
19 optimizer = optim.SGD(model.parameters(), lr=0.001)
```

```
In [24]: 1 import time
2 start_time = time.time()
3
4 for epoch in range(10):
5     total_loss = 0
6     for context, target in data:
7
8         # Step 1. Prepare the inputs to be passed to the model (i.e, turn the words
9         # into indices)
10         context_idxs = torch.tensor([word_to_ix[w] for w in context], dtype=torch.long)
11
12         # Step 2. Recall that torch *accumulates* gradients. Before passing in a
13         # new instance, you need to zero out the gradients from the old instance
14         model.zero_grad()
15
16         # Step 3. Run the forward pass, getting log probabilities over next words
17         log_probs = model(context_idxs)
18
19         # Step 4. Compute your loss function. (Again, Torch wants the target word
20         # index as an input)
21         loss = loss_function(log_probs, torch.tensor([word_to_ix[target]], dtype=torch.long))
22
23         # Step 5. Do the backward pass and update the gradient
24         loss.backward()
25         optimizer.step()
26
27         # Get the Python number from a 1-element Tensor by calling tensor.item()
28         total_loss += loss.item()
29     losses.append(total_loss)
30 print(losses) # The loss decreased every iteration over the training data!
31
32 end_time = time.time()
33 print(f"运行时间: {end_time - start_time} ")
```

[383040.6752166748, 371139.4347035885, 349080.9028342962, 342351.5752040148, 339795.038377285, 337973.7778482735, 336484.98399430513, 335140.70063331723, 333851.7763360292, 332568.11263199896, 331258.9643229209, 329905.11920079216, 328492.62182867713]

运行时间: 820.6160748004913

```
In [25]: 1 # Create your model and train. Here are some functions to help you make the data
2 def make_context_vector(context, word_to_ix):
3     idxs = [word_to_ix[w] for w in context]
4     return torch.tensor(idxs, dtype=torch.long)
5
6 make_context_vector(data[0][0], word_to_ix) # example
```

```
Out[25]: tensor([ 764, 2245,  764, 2245])
```

```
In [ ]: 1
```

使用gensim中的Word2Vec包实现CBOW

```
In [43]: 1 from gensim.models import Word2Vec
```

```
In [ ]: 1 """
2 - sentences: 可迭代的语句列表，较大的语料库可以考虑从磁盘/IO的形式传输
3 - vector_size: 单词向量的维数
4 - window: 句子中当前单词与预测单词的最大距离
5 - min_count: 忽略总频率低于此值的所有单词
6 - workers: 使用多个 worker 线程训练模型
7 - sg: 训练算法，1-> skip-gram 否则 -> CBOW
8 - hs: 1 -> 分层 softmax 方法，否则 -> 负采样
9 - negative: >0 则使用负采样，通常推荐距离为 [5-20]，如果设置为0则不适用负采样
10 - alpha: 初始学习率
11 - min_alpha: 随着训练进行，学习率将线性下降至 min_alpha
12 - max_vocab_size: 词库限制，每 1000w 个字类型大约需要1GB的 RAM
13 - sample: 配置较高频率的单词随机下采样的阈值，生效范围 (0, 1e-5)
14 - epochs 迭代次数
15 """
```

```
In [48]: 1 with open("三国演义.txt", 'r', encoding='utf-8') as f: # 读入文本
2     lines = []
3     for line in f: #分别对每段分词
4         temp = jieba.lcut(line) #结巴分词 精确模式
5         words = []
6         for i in temp:
7             #过滤掉所有的标点符号
8             i = re.sub("[\s+\.!\\/_,$%^*(+\"'\" \" 《》 ]+|[+——! , 。 ? 、 ~@#¥%·",
9             if len(i) > 0:
10                 words.append(i)
11             if len(words) > 0:
12                 lines.append(words)
13 print(lines[0:5])
```

```
[['三国演义'], ['第一回', '宴', '桃园', '豪杰', '三', '结义', '斩', '黄巾', '英雄', '首', '立功'], ['滚滚', '长江', '东', '逝水', '浪花', '淘尽', '英雄', '是非成败', '转头', '空'], ['青山', '依旧', '在', '几度', '夕阳红', '白发', '渔樵', '江渚上', '惯'], ['看', '秋月春风', '一壶', '浊酒', '喜相逢', '古今', '多少', '事', '都', '付']]
```

```
In [60]: 1 model = Word2Vec(
2         lines,
3         vector_size = 100,
4         hs = 1,
5         sg = 2,
6         min_count = 1,
7         window = 2,
8         workers = 4,
9         epochs = 10
10      )
```

```
In [61]: 1 print("孔明的词向量: \n", model.wv.get_vector('孔明'))
```

孔明的词向量:

```
[-0.32365045  0.27591166  0.5305709  -0.23610966 -0.40783113 -0.05334409
-0.12471253  0.30638564  0.75083137 -0.64840454 -0.0165752  0.40281984
 0.03974082 -0.30338097  0.0661559  -0.32427162 -0.24906923 -0.40847516
-0.39091098  0.11982058 -0.14609097 -0.0076288  0.3399963  0.00193367
 0.72057104  0.21825542 -0.9662531  -0.35575494 -0.26800779 -0.22006762
-0.3623526  -0.06418754  0.0697577  -0.2002573  -0.0141564  0.702028
 0.02826011  0.31054723  0.55619836 -0.38314077  0.50334597 -0.4325456
-0.5552436  -0.50480545  0.12148239 -0.66760194 -0.63635737  0.40741077
 0.39663538 -0.09630173 -0.24927041 -0.04164707 -0.17518984  0.7221643
 0.3903633  -0.00678943  0.5938008  -0.26348865 -0.41011506 -0.40510067
-0.39398313 -0.26716083  0.5274327  0.37557262  0.47788078  0.1622746
 0.3886236  0.7362393  -0.31729952  0.41644067 -0.6646033  0.21656819
 0.72482973  0.10144351  0.38675648 -0.02992108 -0.25849965 -0.15729703
-0.58085984  0.19022861 -0.04304889 -0.44869766 -0.03818106  0.08713236
-0.29699644  0.07344805  0.59316033  0.41191667 -0.12846035 -0.03103376
 0.12596759 -0.10565289  0.05203108 -0.3798121  -0.15744796  0.07548639
 0.02230603 -0.34300146  0.421915  -0.4369814 ]
```

```
In [62]: 1 print("\n和孔明相关性最高的前20个词语:")
2 model.wv.most_similar('孔明', topn = 20)# 与孔明最相关的前20个词语
```

和孔明相关性最高的前20个词语:

```
Out[62]: [('玄德', 0.7855352759361267),
('孟获', 0.7491429448127747),
('先主', 0.7122588157653809),
('懿', 0.71099853515625),
('周瑜', 0.7088468074798584),
('鲁肃', 0.7035942673683167),
('姜维', 0.6889432072639465),
('孙夫人', 0.686983048915863),
('郝昭', 0.6833346486091614),
('心中', 0.6751553416252136),
('孔明来', 0.6660807728767395),
('瑜', 0.662392258644104),
('庞统', 0.6611930727958679),
('孙权', 0.6606889367103577),
('后主', 0.6582843661308289),
('关公', 0.6560310125350952),
('王允', 0.6555484533309937),
('马谡', 0.6552057862281799),
('孔明回', 0.6539878845214844),
('玄德公', 0.6519277095794678)]
```

```
In [63]: 1 def getSimilarSeq(key, model, top = 10):
2         print("Top For %s =====" % key)
3         sims = model.wv.most_similar(key, topn = top)
4         for i in sims:
5             print(i)
6         print("End Sim For %s =====" % key)
```

```
In [64]: 1 getSimilarSeq("玄德", model)
2         getSimilarSeq("云长", model)
```

```
Top For 玄德 =====
('孔明', 0.7855352163314819)
('云长', 0.7299472689628601)
('孙夫人', 0.7004892826080322)
('鲁肃', 0.698726236820221)
('孔明遂', 0.6961550712585449)
('孟获', 0.6815526485443115)
('庞统', 0.6662197113037109)
('诸葛均', 0.6655987501144409)
('孙权', 0.6647616624832153)
('孙乾', 0.6642411351203918)
End Sim For 玄德 =====
Top For 云长 =====
('张飞', 0.7685154676437378)
('赵云', 0.7665302753448486)
('关公', 0.7517321109771729)
('岱', 0.7360236644744873)
('玄德', 0.7299474477767944)
('孙夫人', 0.7195055484771729)
('王忠', 0.7171892523765564)
('魏延', 0.7167795896530151)
('翼德', 0.7129519581794739)
('忠', 0.7106291055679321)
End Sim For 云长 =====
```