

Topic 4: Deep Learning (Part I)

Tian Xie[†]

[†]**Shanghai University of Finance and Economics**

November 1, 2023 @ **Techno-Economics Research Alliance (TERA)**

A Nibble of History

- ▶ Deep learning has its theoretical foundation dating back to the 1940s.
- ▶ The field has undergone various name changes throughout its history, including “cybernetics” in the 1940s-1960s and “connectionism” in the 1980s-1990s.
- ▶ The term “deep learning” gained prominence in 2006, marking a resurgence in the field.
- ▶ Deep learning is sometimes confused with artificial neural networks (ANN), as early algorithms aimed to mimic biological learning processes.
- ▶ The modern concept of deep learning extends beyond a neuro-scientific perspective, involving the learning of multiple levels of composition.
- ▶ It is considered a more general principle that surpasses conventional/classic machine learning approaches.

Network Structure

- ▶ Feedforward networks consist of **layers** of interconnected **neurons**.
- ▶ Layers can be categorized into input, hidden, and output layers.
 - ▶ **Graph Representation:** Neural networks use an acyclic graph, like $f(\mathbf{x}) = f^{(3)}(f^{(2)}(f^{(1)}(\mathbf{x})))$, depicting the composition of functions.
 - ▶ **Depth and Output:** The chain's length determines depth in *deep learning*. The last layer, the output layer, matches and produces target outputs during training.
 - ▶ **Hidden Layers:** Situated between input and output, hidden layers are crucial for learning and approximating f^* . The algorithm decides their role for optimal approximation.
 - ▶ **Neuroscience Inspiration:** Neural networks borrow from neuroscience, using vector-valued hidden layers with width determined by dimensionality. Layers parallelly compute, mirroring the brain's structure.

Activation Functions

- ▶ **Non-Linearity Introduction:** Activation functions bring non-linearity, enabling the network to model complex relationships.
- ▶ **Common Choices:** ReLU, sigmoid, and tanh are widely used activation functions.
 - ▶ **Rectified Linear Unit (ReLU):** $f(x) = \max(0, x)$
 - ▶ **Sigmoid:** $f(x) = \frac{1}{1+e^{-x}}$
 - ▶ **Hyperbolic Tangent (tanh):** $f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$
- ▶ **Custom Activation:** Depending on the problem, custom activation functions can be designed.
- ▶ **Activation Impact:** The choice influences the network's learning capacity and convergence speed.

Cost Function

- ▶ **Task Dependence:** The choice of cost function is task-specific.
- ▶ **Classification:** Cross-entropy is commonly used for classification tasks.
- ▶ **Regression:** Mean Squared Error (MSE) or Mean Absolute Error (MAE) are preferred for regression.
- ▶ **Customization:** Depending on the problem, custom cost functions can be designed.
- ▶ **Loss Landscape:** The cost function shapes the optimization landscape during training.
- ▶ **Trade-offs:** Selection involves trade-offs between robustness, sensitivity, and computational efficiency.

Gradient-Based Learning

- ▶ **Weight Tuning:** Feedforward network training involves adjusting weights to minimize the chosen cost function.
- ▶ **Non-Convex Challenges:** The cost function's non-convex nature poses optimization difficulties during training.
- ▶ **Local Minima:** Networks may converge to suboptimal local minima due to non-convexity.
- ▶ **Gradient Descent:** Gradient-based methods like backpropagation are commonly used for weight updates.
 - ▶ **Momentum:** GD updates parameters based only on the current gradient. Momentum considers both the current gradient and a historical moving average of gradients, improving convergence and reducing oscillations.
 - ▶ **Exploration Strategies:** Techniques like stochastic GD and adaptive learning rates address optimization challenges.

Forward Propagation and Backprop

- ▶ **Forward Propagation:** Input passes through layers, computing weighted sums, applying activations, and passing information to subsequent layers.
- ▶ **Backprop:** Backward propagation of errors, the core training algorithm for neural networks, involves:
 - ▶ **Error Calculation:** Compare predictions to target values using the chosen cost function.
 - ▶ **Gradient Descent:** Propagate error backward, computing gradients of the cost function with respect to weights and biases.
 - ▶ **Weight Updates:** Update weights and biases using computed gradients and an optimization algorithm, typically gradient descent, to minimize error and improve model performance.

Regularization Techniques

- ▶ **Norm Penalties:** Regularization terms are added to the loss function, penalizing large weights. L1 regularization enforces sparsity, while L2 regularization discourages overly complex models.
- ▶ **Early Stopping:** Training is halted when the model's performance on a validation set stops improving, preventing overfitting to the training data.
- ▶ **Dropout:** Randomly deactivates a fraction of neurons during training, promoting robustness and preventing reliance on specific neurons.
- ▶ **Data Augmentation:** Introduces variations in training data by applying transformations like rotation, scaling, or flipping, enhancing model generalization.
- ▶ **Batch Normalization:** Normalizes input data within each mini-batch, reducing internal covariate shift and improving convergence.

Convolution Neural Network

Convolution Operation

► Convolution Overview

- The convolution operation is the cornerstone of CNNs, involving the application of a filter (kernel) to the input data.
- Mathematically, convolution is expressed as:

$$(x * w)(t) = \int_{-\infty}^{\infty} x(a) \cdot w(t - a) da$$

where x is the input and w is the filter.

► Feature Extraction

- Convolution enables the network to learn hierarchical features by capturing local dependencies within the input data.
- This process is essential for identifying patterns and structures in the data.

Motivation for Convolution

▶ **Sparse Interactions:**

- ▶ Sparsity reduces memory requirements and computation, enhancing efficiency.
- ▶ Allows capturing meaningful features with smaller kernels.

▶ **Parameter Sharing:**

- ▶ Sharing parameters across input positions dramatically improves efficiency.
- ▶ Convolution is equivariant to translation, providing consistent representations.

▶ **Equivariant Representations:**

- ▶ Convolution exhibits equivariance to translation, useful for creating timelines in time series data.
- ▶ Not naturally equivariant to scale or rotation, requiring additional mechanisms.

Pooling

- ▶ Convolutional layers operate in three stages:
 - ▶ Parallel convolutions produce linear activations.
 - ▶ Nonlinear activation functions (e.g., ReLU) in the detector stage.
 - ▶ Pooling modifies the output using summary statistics.
- ▶ Pooling achieves **invariance to small translations**.
 - ▶ Adds an **infinitely strong prior** for translation invariance.
 - ▶ Reduces computational and memory requirements.
 - ▶ Essential for handling inputs of varying size.
- ▶ Different pooling functions exist (e.g., max pooling, average pooling).
- ▶ Dynamic pooling and theoretical guidance for various situations.

Convolution & Pooling as Strong Priors

- ▶ Priors can be weak or strong based on probability density concentration.
- ▶ **Infinitely strong priors** strictly restrict some parameters.
- ▶ Convolution as a fully connected net with a strong prior:
- ▶ Promotes local interactions and equivariance to translation.
- ▶ Pooling as an **infinitely strong prior** for translation invariance.
- ▶ Can cause **underfitting** if assumptions are inaccurate.
- ▶ Models designed for balancing translation invariance and information preservation.
- ▶ Compare convolutional models to benchmarks of statistical learning performance.

Stride and Padding

► **Stride:**

- Controls the step size of the filter during convolution.
- Larger strides result in smaller output spatial dimensions.
- Influences overlap, computational efficiency, and information preservation.

► **Padding:**

- Padding adds extra border pixels to the input before convolution.
 - Mitigates spatial dimension reduction, preserving input information.
 - Common padding types include zero-padding.
- Which one is more useful in time series analysis?